



**Faculty of Electrical Engineering
Department of Computer Science**

Bachelor's Thesis

Web-based application for managing private bank accounts

Yelyzaveta Taranova

Software Engineering and Technology

May 2026

Supervisor: Ing. Ondřej Vondrouš, Ph.D.

Acknowledgement / Declaration

I would like to express my gratitude to my supervisor, Ing. Ondřej Vondrouš, Ph.D., for his guidance during the creation of this bachelor's thesis.

I am also thankful to my work colleagues. Inspiring environment they provided helped me greatly in balancing my professional duties with my academic work.

Finally, I would like to thank my friends and family for their support and motivation throughout my studies.

I declare that I have prepared the submitted thesis independently and that I have cited all the information sources used in accordance with the Methodological Guideline on the observance of ethical principles in the preparation of university final theses.

In Prague 30.04.2026

Abstrakt / Abstract

Tento projekt se zaměřuje na návrh a vývoj webové aplikace pro správu osobních bankovních účtů prostřednictvím agregace finančních údajů z více českých bank na jednom místě.

Úvod nejprve shrnuje hlavní motivaci a základní cíle projektu. Následující část se věnuje analýze existujících tržních alternativ za účelem stanovení požadavků a očekávání pro nový systém. Poté následuje kapitola o návrhu, která definuje celkovou architekturu, strukturu databáze a externí spojení potřebná k vytvoření platformy.

Kapitola o implementaci vysvětluje samotný proces vývoje a funkce aplikace. Kapitola o testování poté popisuje strategie použité k ověření, že software funguje správně. Závěr nakonec hodnotí aktuální stav projektu a navrhuje možná budoucí vylepšení pro rozšíření jeho funkcí.

Klíčová slova: Bankovní aplikace, osobní bankovníctví, centralizace financí, webová aplikace, Spring Boot Framework, REST API, Java, SOA.

This project focuses on the design and development of a web application for managing personal bank accounts by aggregating financial data from multiple Czech banks in one place.

First, the introduction outlines the core motivation and the primary goals of the project. Next, the analysis section reviews current market alternatives to establish a set of requirements and expectations for the new system. The design chapter follows, defining the overall architecture, database layout, and external connections required to build the platform.

The implementation chapter explains the actual development process and application functionalities. The testing chapter then outlines the strategies used to verify that the software functions correctly. Finally, the conclusion evaluates the actual application state and suggests potential future updates to expand its capabilities.

Keywords: Banking application, personal banking, finance centralization, web application, Spring Boot Framework, REST API, Java, Service-Oriented Architecture.

Contents /

1 Introduction	1	
1.1 Motivation	1	
1.2 Objectives	1	
2 Existing Solutions Overview	3	
2.1 Budgeting Applications	3	
2.2 Wealth Trackers	3	
2.3 Global Solutions with European Support	4	
2.4 Summary	4	
2.5 Key Features in Existing Finance Applications	4	
3 Analysis	6	
3.1 Czech Banking APIs Analysis	6	
3.1.1 The Payment Services Directive	6	
3.1.2 Authentication and Security Protocols	6	
3.1.3 Sandbox vs. Production Environments	7	
3.1.4 Comparison of Czech Banking APIs	7	
3.1.5 Conclusion	7	
3.2 Business Analysis	8	
3.2.1 Functional Requirements	8	
3.2.2 Non-Functional Requirements	11	
3.2.3 Use cases	11	
3.2.4 Process Analysis	14	
3.3 Technology Analysis	16	
3.3.1 Backend Development	16	
3.3.2 Frontend Development	17	
4 Proposed Solution	19	
4.1 Architecture	19	
4.1.1 Architectural Style	19	
4.1.2 Architecture Diagram	20	
4.1.3 Internal Architecture	20	
4.1.4 Component Diagram	21	
4.2 Data design	23	
4.2.1 Database Model	23	
4.2.2 Domain Class Model	26	
5 Implementation	30	
5.1 Features	30	
5.1.1 Dashboard	30	
5.1.2 Banks Overview	31	
5.1.3 Accounts Overview	33	
5.1.4 Transactions Overview	35	
5.1.5 Budgeting	36	
5.1.6 Categorization Service	39	
5.1.7 Currency Exchange	42	
5.2 Frontend Implementation	42	
5.3 Security	43	
5.3.1 Keycloak	43	
5.3.2 Authorization Service	44	
5.4 Caching and Resilience	47	
5.5 Deployment	51	
5.5.1 Continuous Integration Pipeline	51	
5.5.2 Docker	51	
5.6 Error Handling	52	
6 Testing	54	
6.1 Unit Testing	54	
6.2 Integration Testing	55	
6.3 End-to-end Testing	57	
7 Conclusion	60	
References	61	
A Zadání práce	65	
B Acronyms	66	

Tables / Figures

3.1 Functional Requirements — Access to Public Pages.....	3.1 Use Cases — Anonymous User.....
8	12
3.2 Functional Requirements — Profile Management.....	3.2 Use Cases — Logged-in User ..
8	13
3.3 Functional Requirements — Bank Integration	3.3 Process Diagram — Connecting a Bank Account
9	15
3.4 Functional Requirements — Financial Data Retrieval.....	4.1 Application Architecture.....
9	20
3.5 Functional Requirements — Visualization	4.2 Internal Architecture.....
10	21
3.6 Functional Requirements — Categorization	4.3 Component Diagram.....
10	22
3.7 Functional Requirements — Transactions Grouping.....	4.4 Authorized Client Table
10	24
3.8 Functional Requirements — Product Filtering.....	4.5 Bank Connection Table
10	24
3.9 Functional Requirements — Budgeting	4.6 Budget Table
11	25
3.10 Functional Requirements — System monitoring	4.7 Categorization Tables
11	25
3.11 Non-functional Requirements..	4.8 Domain Class Model — User Service
11	26
3.12 Selected technologies for application backend development	4.9 Domain Class Model — Bank Connection Service
17	26
6.1 End-to-End Testing Scenarios .	4.10 Domain Class Model — Product Service
59	27
	4.11 Domain Class Model — Shared Service.....
	28
	5.1 Feature — Dashboard
	30
	5.2 Feature — Banks Overview....
	31
	5.3 Feature — Accounts Overview
	33
	5.4 Feature — Transactions Overview
	36
	5.5 Feature — Budgeting
	37
	5.6 Keycloak Administrator Dashboard
	44
	5.7 Circuit Breaker Pattern.....
	50
	5.8 Passed CI Pipeline
	51

Chapter 1

Introduction

Keeping track of personal finances is getting harder since most people now use multiple bank accounts and cards from different institutions. Czech banks generally have good online platforms, but checking your overall financial situation usually means logging into several different apps to piece the information together.

Thanks to PSD2¹ and open banking rules in the European Union, it is now possible for third-party apps to access bank data securely, as long as the user gives permission. This thesis takes advantage of that to build a unified personal finance dashboard.

The goal of this project is to develop a web application that uses these modern banking APIs² to pull a user's accounts into one place. Through a single interface, users can check their current balances, track transactions across their various Czech bank accounts, and manage their budgets. The application is designed purely for personal use and focuses on organizing data safely. To keep the project secure and focused, it only reads data and strictly avoids active operations like sending payments or changing account settings.

1.1 Motivation

Even though Czech banks support PSD2 APIs, their systems are still isolated from one another. If a user has accounts across Česká spořitelna, Air Bank, and Moneta, they have to jump between different apps just to see how much money they actually have.

There are international apps that try to solve this, but they usually struggle to connect with Czech banks. When they do work, they often rely on third-party data aggregators, which raises privacy concerns and usually means the user's data is passing through extra hands.

Because of this, the project is motivated by four main goals:

- **Centralization** — pulling all accounts and transaction history into a single application with an easy-to-read dashboard.
- **Transparency** — giving the user a clear, instant view of their cash flow without requiring them to do manual calculations.
- **Control and privacy** — keeping the user's financial data secure and private without relying on third-party data brokers.
- **Simplicity** — building a straightforward personal tool instead of a massive, complicated fintech product.

1.2 Objectives

The main goal of this thesis is to build a web application that safely connects to different Czech banks and helps the user make sense of their money. Specifically, the project focuses on the main tasks:

¹ PSD2 — Payment Services Directive 2

² API — Application Programming Interface

2. Build a single dashboard for all accounts.

3. Make spending habits easy to understand.

4. Put open banking to practical use.

5. Learn from existing applications.

2

Chapter 2

Existing Solutions Overview

Before designing the new application, it is important to analyze how existing tools handle personal finance management. The core features across the industry are generally the same: pulling transaction data, categorizing spending, and providing visual budgeting tools. Applications differ significantly in their technical implementation, specifically in how they access banking data and their geographic focus. This section evaluates several popular solutions to identify their strengths, design patterns, and limitations regarding Czech banking systems.

2.1 Budgeting Applications

Applications built for the US market often set the standard for user interface design and budgeting methodology.

- **You Need A Budget (YNAB)** uses a strict budgeting approach (zero-based budgeting) where every unit of income is assigned to a specific category before it is spent. It simplifies spending decisions and clarifies priorities so users can stop worrying about money.[1]
- **Monarch Money** focuses on household finances, allowing multiple users to manage a shared database. It also actively scans transaction histories to flag recurring subscriptions. [2]
- **PocketGuard** calculates the user's exact disposable income. The application subtracts scheduled upcoming expenses from the current account balance, providing a direct metric of the funds currently available for discretionary spending. [3]

Limitations: although functionally advanced, these applications share a major technical flaw for European users. They rely almost exclusively on North American API aggregators like Plaid. Because they lack integration with the European PSD2 framework, they cannot connect to Czech banks.

2.2 Wealth Trackers

- **Empower** is designed to track entire financial portfolios. It aggregates daily checking accounts alongside investment portfolios, retirement funds, and real estate values. The application's primary function is to calculate total net worth and analyze long-term asset allocation. [4]
- **Quicken Simplifi** connects to bank accounts, credit cards, and investment accounts to provide real-time updates on cash flow and spending trends. Users can create custom budgets, establish savings goals, and track specific spending categories through personalized watchlists.[5]

Limitations: tracking multiple asset classes requires a complex user interface. Both applications share the same geographic data restrictions. They rely entirely on North American financial aggregators and lack direct API support for Central European financial institutions.

2.3 Global Solutions with European Support

- **Spendee** is an application with Czech origins that maintains a large global user base. It functions properly within the Czech Republic by offering detailed transaction categorization and handling multiple currencies. It can automatically synchronize with local institutions, providing an automated overview of daily finances for regional users. Spendee was the first to obtain a PSD2 license in the Czech Republic. [6]
- **BudgetBakers** is another globally recognized personal finance application developed in the Czech Republic. Similar to Spendee, it provides advanced transaction categorization, budgeting tools, and multi-currency support. It also features automatic data synchronization with Czech banks. [7]

Limitations: to support thousands of banks worldwide, neither application connects to financial institutions directly. Instead, they rely on third-party data aggregators like Salt Edge. This approach introduces a middleman into the data flow. The user must trust a third-party broker with their financial consent and data, which creates potential privacy and security concerns compared to direct bank connections.

2.4 Summary

Analyzing these existing solutions highlights a clear gap that this thesis aims to fill. The applications with the best interfaces (YNAB, Monarch) cannot connect to Czech banks. The applications that can connect to local banks rely on third-party data brokers, adding a layer of privacy risk.

This justifies the core technical approach of this project: building a straightforward, localized tool that connects directly to Czech bank APIs. By skipping the third-party aggregators, the application can guaranty that the user's data travels strictly between their bank and their own private dashboard.

2.5 Key Features in Existing Finance Applications

Based on the analysis of existing solutions, standard personal finance applications typically implement a common set of core functionalities:

- **Account Aggregation** — Connecting various financial accounts to consolidate balances and transaction histories into a single interface. The standard approach involves automated, periodic data synchronization via APIs or data brokers.
- **Transaction Categorization** — Automatically importing and assigning metadata (categories such as Groceries or Utilities) to individual transactions.
- **Budget Management** — Creating spending limits per category on a defined cycle, usually monthly.
- **Goal Tracking** — Defining specific financial targets, such as accumulating an emergency fund or reducing specific debts.
- **Recurring Payment Detection** — Analyzing transaction histories to identify predictable, repeating charges, like utility bills or software subscriptions. This feature allows users to anticipate upcoming cash flow requirements and identify redundant services.
- **Data Visualization** — Generating charts and graphs to visually represent financial states. Standard outputs include charts for categorical spending, graphs for cash flow trends, and historical net worth progression.

- **Automated Notifications** — Triggering system alerts based on specific account states, such as approaching a budget limit, identifying a critically low balance, or detecting unusual transaction patterns.
- **Security and Synchronization** — Implementing strong encryption and multi-factor authentication to protect sensitive financial data.

Chapter 3

Analysis

This chapter examines the project requirements and the chosen technical solutions. It is divided into three main parts.

First, the external API analysis explores the integration with third-party banking interfaces, explaining how the application connects to testing sandboxes to retrieve financial data safely. Second, the business analysis defines the core goals of the application and analyzes the necessary system features. Finally, the technology analysis evaluates the tools and frameworks selected to build the frontend and backend.

3.1 Czech Banking APIs Analysis

Modern personal finance applications require secure integrations with banking systems. In the Czech Republic, major banks expose standardized APIs mandated by European regulations. These interfaces allow third-party providers to access account information following explicit user consent. The APIs generally align with the Czech Open Banking Standard or the Berlin Group framework, providing consistent data structures for retrieving account details, balances, and transaction histories. This section outlines the regulatory foundation, technical capabilities, security requirements, and access conditions for these APIs.

3.1.1 The Payment Services Directive

The revised Payment Services Directive (PSD2) is a European Union regulatory framework that fundamentally changed how financial data is managed. Prior to this legislation, banks held exclusive control over their customers' account information. PSD2 mandates that all European financial institutions must provide secure, standardized API access to licensed Third-Party Providers.

This regulation established two primary types of access: Payment Initiation Services (PIS) and Account Information Services (AIS). For the scope of this application, only AIS is relevant. Under the AIS framework, a user can grant a third-party application explicit consent to retrieve their account information and transaction history in a read-only format. The directive ensures that users maintain complete ownership of their financial data and strictly control who can access it. [8]

3.1.2 Authentication and Security Protocols

Authentication relies strictly on the OAuth ¹ 2.0 protocol combined with Strong Customer Authentication (SCA). During the synchronization process, the application redirects the user to the bank's own secure portal. At this stage, the user must complete the SCA process to formally grant access. This authorization flow ensures that the custom application never handles or stores direct login credentials. All subsequent data exchanges require mutual TLS ² encryption to ensure transport security. [9]

¹ OAuth — Open Authorization

² TLS — Transport Layer Security

■ 3.1.3 Sandbox vs. Production Environments

Accessing banking APIs depends entirely on the deployment environment. Connecting to production APIs to fetch real user data requires a valid eIDAS ³ certificate and an official license from the Czech National Bank to operate as a regulated Account Information Service Provider. Acquiring this regulatory approval is highly complex and generally not feasible for private tools or academic projects.

To resolve this access limitation, financial institutions provide developer sandbox environments. These sandboxes replicate the exact architecture and endpoint structure of the production APIs and return realistic, mock financial data. Developers can register applications, implement the OAuth authentication flow, and test data retrieval processes without holding a formal PSD2 license.

■ 3.1.4 Comparison of Czech Banking APIs

To successfully integrate multiple financial institutions into a unified application, it is necessary to evaluate the technical capabilities and access requirements of their respective APIs. This section provides a comparative analysis of the major banks used in the Czech Republic: Komerční banka [10], Česká spořitelna [11], Fio Bank [12], Air Bank [13], Raiffeisenbank [14], ČSOB [15], Moneta Bank [16], and UniCredit [17].

3.1.4.1 Provided Features

All mentioned financial institutions fully support the core Account Information Services required for a personal finance dashboard. Specifically, their APIs provide standardized endpoints for **listing personal accounts**, **listing current account balances**, and **listing transaction history**.

3.1.4.2 Production Environment Requirements

To access the production environments and retrieve real user data, all six institutions enforce strict regulatory and technical prerequisites. Every bank requires the applicant to hold a valid PSD2 AISP ⁴ license. From a technical perspective, establishing a production connection universally requires a Qualified Website Authentication Certificate (QWAC).

3.1.4.3 Sandbox Environment Requirements

The requirements for sandbox environments are significantly less restrictive. For all of the evaluated financial institutions, developers simply need to complete a standard registration process on their respective developer portals to obtain sandbox credentials, API keys, and test certificates.

All mentioned banks, except Moneta Money Bank, Fio Bank, and UniCredit, allow independent developers and private applications to access their testing infrastructure. Moneta Money Bank, Fio Bank, and UniCredit Bank prohibit academic projects and private applications from using their testing platforms, making their APIs unavailable for independent development and research.

■ 3.1.5 Conclusion

The analysis of banking APIs available in the Czech Republic shows that major banks offer highly standardized Account Information Services driven by the PSD2 framework.

³ eIDAS — electronic IDentification, Authentication and trust Services

⁴ AISP — Account Information Service Providers

Access to real production data remains strictly regulated and limited to licensed third-party providers. The sandbox environments offered by these banks significantly lower the access barriers. These simulated environments provide a highly suitable platform for development, prototyping, and fully verifying functional architecture. This project will make use of the accessible sandbox APIs: Komerční banka, Česká spořitelna, Air Bank, Raiffeisenbank and ČSOB.

3.2 Business Analysis

This section defines the business context of the application and translates user needs into structured functional and non-functional requirements. It identifies the target audience, describes expected system behavior, and outlines key business processes supported by the application. The goal of this analysis is to ensure that the proposed solution aligns with real user needs while remaining technically feasible and compliant with regulatory constraints.

3.2.1 Functional Requirements

This section specifies the functional requirements of the application, defining what the system shall do from the user's perspective. The requirements are grouped by functional areas and cover all major interactions, including authentication, bank integration, financial data retrieval, visualization, budgeting, and administration.

3.2.1.1 Access to Public Pages

ID	Title	Description
FR1	Displaying the landing page	The system shall allow anonymous users to view the landing page with a general description of the application and its purpose, information about features, security, and supported banks.

Table 3.1. Functional Requirements — Access to Public Pages

3.2.1.2 Profile Management

ID	Title	Description
FR2	Registration	The system shall allow anonymous users to create an account via the registration form.
FR3	Log in	The system shall support logging in using an email and password.
FR4	Log out	The system shall allow the user to log out.

Table 3.2. Functional Requirements — Profile Management

3.2.1.3 Bank Integration

ID	Title	Description
FR5	Linking a bank	The system shall allow users to link a bank from defined list of financial institutions
FR6	Token storage	The system shall securely store tokens required for retrieving data.
FR7	Bank disconnection	The system shall allow users to disconnect any previously linked bank.

Table 3.3. Functional Requirements — Bank Integration

3.2.1.4 Financial Data Retrieval

ID	Title	Description
FR8	Product retrieval	The system shall retrieve customers products from all linked banks.
FR9	Account balances retrieval	The system shall retrieve the balance for each account from all linked banks.
FR10	Transaction history retrieval	The system shall retrieve the transaction history for each account from all linked banks for the specified time period.
FR11	Financial data caching	The system shall cache financial data periodically.

Table 3.4. Functional Requirements — Financial Data Retrieval

3.2.1.5 Visualization

ID	Title	Description
FR12	Displaying a dashboard	The system shall display a consolidated dashboard containing an overview of assets and liabilities.
FR13	Displaying statistics	The system shall present graphs of spending and income for the selected time period.
FR14	Displaying transactions	The system shall display transactions in a clear list format for every account from all linked banks.
FR15	Displaying products	The system shall display products with balances in a clear list format for every account from all linked banks.
FR16	Displaying budgets	The system shall display budgets in a clear list format.
FR17	Displaying profile	The system shall user profile.

Table 3.5. Functional Requirements — Visualization

3.2.1.6 Categorization

ID	Title	Description
FR18	Transactions categorization	The system shall automatically categorize transactions using predefined rules.
FR19	Products categorization	The system shall automatically categorize products using predefined rules.

Table 3.6. Functional Requirements — Categorization

3.2.1.7 Transactions Grouping

ID	Title	Description
FR20	Transactions Grouping	The system shall group transactions by month and by category.

Table 3.7. Functional Requirements — Transaction Grouping

3.2.1.8 Product Filtering

ID	Title	Description
FR21	Product Filtering	The system shall allow the filtering of products by bank.

Table 3.8. Functional Requirements — Product Filtering

3.2.1.9 Budgeting

ID	Title	Description
FR22	Setting monthly budgets	The system shall allow users to set monthly budgets for categories.
FR23	Budget warnings	The system shall display warnings or highlights when a budget is exceeded.

Table 3.9. Functional Requirements — Budgeting

3.2.1.10 System Monitoring

ID	Title	Description
FR24	System logs	The system shall provide logs for authentication attempts, API errors, and system faults.

Table 3.10. Functional Requirements — System monitoring

3.2.2 Non-Functional Requirements

ID	Title	Description
NFR1	Security	The application must use secure storage and must not store sensitive data in an open form.
NFR2	Resilient system	The system must handle API failures gracefully.
NFR3	Application performance	The application should load the dashboard within 20 seconds.
NFR4	Browser compatibility	The application must function properly on the current versions of major web browsers.

Table 3.11. Non-functional Requirements

3.2.3 Use cases

This section presents the main use cases of the system, describing how different types of users interact with the application. Use cases are divided according to user roles and illustrate the scope of functionality available to each role. The diagrams provide a high-level overview of system behavior and support the understanding of functional requirements.

3.2.3.1 Anonymous User

Anonymous users are visitors who access the application without creating an account or logging in.




Figure 3.1. Use Cases — Anonymous User

This diagram illustrates the initial interaction points available to an unauthenticated individual, designated as the **Anonymous User** actor. Within the system boundary of the web application, this actor has direct associations with three primary use cases: viewing the public-facing landing page, creating a new account, and authenticating into the system.

These three use cases directly implement the foundational access and authentication functional requirements defined in section 3.2.1:

- **View landing page** — This use case implements FR1 (Displaying the landing page) from section 3.2.1.1. It ensures the system allows unauthenticated users to read general information regarding the application’s purpose, features, and security.
- **Register** — This use case implements FR2 (Registration) from section 3.2.1.2. It represents the workflow where an anonymous user submits their credentials via a registration form to create a new profile and gain system access.
- **Log in** — This use case implements FR3 (Log in) from section 3.2.1.2. It fulfills the requirement for the system to authenticate a user via their credentials, successfully transitioning their state from an Anonymous User to a Logged-in User.

3.2.3.2 Logged-in User

A logged-in user has access to the core features of the application, including connecting banking accounts and viewing financial data.




Figure 3.2. Use Cases — Logged-in User

This diagram illustrates the set of interaction scenarios available to a user who has successfully authenticated into the system, designated as the **Logged-in User** actor. Within the web application boundary, this actor maintains direct associations with five primary use cases. These base use cases are further modified through the specific extension relationship **<<extend>>** to encompass the full suite of functional requirements.

The primary use cases and their corresponding extensions implement the system's core capabilities:

- **View financial data** — This use case serves as the central hub for financial visualization and implements functional requirements from section 3.2.1.5. It is extended by specialized actions allowing the user to view the dashboard (FR12), browse transactions (FR14), browse products (FR15), browse linked banks, and filter products (FR21).
- **View profile** — This use case implements FR17 from section 3.2.1.5, enabling the user to access and review their personal account information.

- **Log out** — This use case implements FR4 from section 3.2.1.2, fulfilling the requirement to securely terminate the active user session.
- **Manage linked banks** — This use case represents the primary interface for banking integration and implements functional requirements from section 3.2.1.3. It is extended by specific actions to link a bank (FR5) and disconnect a bank (FR7).
- **Manage budgets** — This use case represents the financial planning functionalities and implements functional requirements from section 3.2.1.9. It is extended by specific actions to set the monthly budget (FR22) and remove the monthly budget.

■ 3.2.4 Process Analysis

Process analysis describes the key workflows supported by the application and explains how individual components interact during typical operations. These processes focus on user interactions with external banking systems, internal data processing, and system responses. The aim is to clarify how functional requirements are realized in practice and to identify critical integration and security steps.

3.2.4.1 Connecting a Bank Account

The process of connecting a bank account enables the user to grant the application access to their financial data through an external banking API.




Figure 3.3. Process Diagram — Connecting a Bank Account

The process begins when a logged-in user selects a bank from the list of supported institutions within the application. After initiating the connection, the backend redirects the user to the selected bank's authentication page. At this stage, the user authenticates directly with their bank using the bank's own security mechanisms, such as two-factor authentication or biometric verification. The application itself never accesses or stores the user's banking credentials.

Once authentication is completed, the bank presents a consent screen where the user explicitly approves access to specific data scopes, typically including account information, balances, and transaction history. After consent is granted, the bank redirects the user back to the application along with an authorization code.

The backend service then exchanges this authorization code for an access token using a secure OAuth2 communication channel. The obtained tokens are securely stored in encrypted form and associated with the user's bank connection. The application uses these tokens to retrieve account and transaction data via the bank's API.

Following successful authorization, the system performs an initial synchronization during which it retrieves the list of available accounts and their associated data. Retrieved accounts are normalized into the internal domain model and cached. Accounts become visible to the user in the application dashboard.

3.3 Technology Analysis

This section analyzes the technologies selected for the development of the web application. The goal is to justify the chosen technology stack based on performance, maintainability, scalability, integration capabilities, and long-term sustainability.

3.3.1 Backend Development

3.3.1.1 Framework

The **Spring Boot** framework offers standardized approaches to dependency injection, configuration management, and modular application design. While lightweight frameworks like Express.js or Flask offer faster initial setup times, Spring Boot excels in complex enterprise environments.

Building upon the robust foundation of the core Java Spring Framework, Spring Boot accelerates the development process by mitigating the need for extensive manual setup.

Spring Boot provides pre-configured dependency bundles known as Spring Starters. For instance, adding the Spring Web or Spring Security starters automatically provisions the project with the necessary tools for web routing and access control.

Spring Boot simplifies deployment by embedding web servers, such as Apache Tomcat, directly into the application artifact. This enables the creation of stand-alone applications that run out of the box, eliminating the traditional requirement of configuring and maintaining external web server environments. [18]

3.3.1.2 Programming Language

Java was chosen as the programming language for backend development. Unlike loosely typed languages, Java provides strong static typing, runs on the highly optimized Java Virtual Machine, and offers enterprise-grade libraries. Java is a high-level, object-oriented, and stable programming language. Its built-in Garbage Collector ensures long-term system stability by automatically managing memory and preventing leaks. [19]

3.3.1.3 Database

PostgreSQL was selected as the relational database management system. A personal finance application requires strict structural consistency. PostgreSQL provides the necessary relational integrity to ensure that financial data remains perfectly organized during complex multi-table operations. [20]

3.3.1.4 Caching

Redis serves as the in-memory data store to accelerate access to frequently requested information. Relying solely on live API calls creates an unacceptable user experience. Compared to simpler local caching solutions like Memcached, Redis supports complex

data structures and allows for highly precise control over cache expiration policies. Storing cached values reduces external API calls and minimizes database read operations. [21]

3.3.1.5 Resilience

Integrating multiple third-party banking APIs introduces a high risk of network failures. Standard error handling is insufficient for external service outages. **Resilience4j** is integrated to implement robust fault-tolerance mechanisms, replacing older libraries like Netflix Hystrix. It provides essential patterns such as circuit breakers, automated retries, and rate limiting. By isolating failing external services and preventing cascading system failures, Resilience4j ensures the primary application remains responsive during temporary outages of specific bank providers. [22]

3.3.1.6 Messaging Configuration

Feign client (Spring Cloud OpenFeign) ensures communication with external banking APIs. Interacting with external endpoints traditionally requires writing repetitive boilerplate code using standard HTTP clients like RestTemplate or WebClient. Feign resolves this complexity by acting as a declarative web service client, which means that the developer must specify the client specification as an interface, and Spring Boot will take care of the implementation. This approach significantly simplifies the codebase, making it highly maintainable when connecting to numerous financial institutions. [23]

3.3.1.7 Conclusion

The selected backend stack establishes a secure, reliable, and highly scalable architecture that is perfectly suited for the strict demands of a financial application. In addition to their proven technical capabilities, these specific technologies were chosen due to the author's prior practical experience with them.

Selected technologies are summarized in the following table.

Category	Technology
Framework	Spring Boot
Programming Language	Java
Database	PostgreSQL
Caching	Redis
Resilience	Resilience4j
Messaging Configuration	Feign Client

Table 3.12. Selected technologies for application backend development

3.3.2 Frontend Development

3.3.2.1 Framework

React is a JavaScript library for building user interfaces, created and maintained by Meta (Facebook). It is currently one of the most popular and widely adopted frontend technologies worldwide. React organizes the interface into reusable components (buttons, forms, graphs, cards, charts). This is ideal for finance dashboards, which require modular widgets. Additionally, React uses a virtual DOM ⁵, which makes UI ⁶ updates

⁵ DOM — Document Object Model

⁶ UI — User Interface

fast when data changes frequently. It has a massive ecosystem of libraries for routing, state management, charts, authentication, and UI styling, which dramatically reduces implementation effort. [24]

Chapter 4

Proposed Solution

This chapter presents the proposed solution by detailing the overall system structure, architectural decisions, and internal organization of the application. It outlines how the individual services are arranged to address the identified business requirements, how data flows securely through the system, and how external banking APIs are integrated. The proposed architecture emphasizes modularity, scalability, and maintainability while strictly adhering to modern security and regulatory standards.

4.1 Architecture

This chapter presents the architectural design of the application, serving as the foundational blueprint for the entire system. Establishing a robust and well-defined architecture is crucial for ensuring the application's scalability, maintainability, and security, especially given the sensitive nature of financial data and the requirement to integrate securely with multiple third-party banking APIs. The following sections detail the chosen architectural style, the high-level system components, the internal structuring of individual services, and the specific boundaries that separate internal logic from external integrations.

4.1.1 Architectural Style

The development of the application is based on Service-Oriented Architecture (SOA). SOA is an architectural style in which a system is built as a collection of modular, reusable, and network-accessible services. Each service represents a distinct business capability — such as user management, bank integration, data synchronization, or reporting — and exposes its functionality through standardized interfaces, typically REST ¹ endpoints. [25]

Although SOA and microservices share conceptual similarities — such as modularization, service boundaries, and network-based communication — they differ significantly in granularity, autonomy, communication style, governance, and deployment. In SOA, services can be independent, but they often share centralized components, such as databases, message busses, or identity providers. In microservices, each service is fully autonomous, owning its own database, technology stack, and deployment pipeline. This eliminates shared dependencies but increases operational complexity. [26]

For this application, SOA provides an optimal balance. It allows for logical separation of domains without the heavy operational overhead of managing distributed databases and complex deployment pipelines required by pure microservices. As the system evolves and the need for higher scalability or independent deployment grows, this SOA foundation can naturally be refined into a microservices architecture. [27]

¹ REST — Representational State Transfer

4.1.2 Architecture Diagram

The architecture diagram illustrates the high-level structure of the application, showing the interaction between the frontend, backend services, external banking APIs, and the database.




Figure 4.1. Application Architecture

The diagram in Figure 4.1 illustrates the Service-Oriented Architecture adopted for the system. The **Frontend** acts as the client layer, initiating requests to the backend. The **Backend** is decomposed into five distinct functional services: the **Authorization service**, the **Bank service**, the **Party service**, the **Product service**, and the **Categorization service**. Unlike a pure microservices architecture, where each service strictly owns its own database, these services share a central **Database** for persistence, ensuring data consistency while keeping the infrastructure manageable. The **Product service** also acts as a gateway to **External APIs**, handling the retrieval of banking data.

4.1.3 Internal Architecture

To ensure separation of concerns and maintainability, each service in the system follows a strict layered architecture composed of three main layers:

1. **Presentation Layer** — this layer exposes REST endpoints that allow clients to interact with the system. It handles incoming HTTP ² requests, validates input structures, and delegates execution to the domain layer. The presentation layer is strictly stateless and does not contain any business logic.
2. **Domain Layer** — this layer encapsulates the core business logic and workflows. It orchestrates operations between repositories, external providers, and domain objects. It is responsible for authorization checks, data transformation, validation rules, and transaction handling. It also contains domain entities and value objects representing the system's primary business concepts.

² HTTP — HyperText Transfer Protocol

3. **Messaging and Persistence Layer** — this layer forms the boundary between the internal business logic and external infrastructure. The Messaging component acts as an adapter for external API implementations. Since each bank may provide different endpoints, data formats, authentication flows, or rate limits, it is essential to centralize and standardize these interactions here.

In addition to external messaging, this bottom layer manages data persistence using JPA ³ and Hibernate. It abstracts database operations and provides a clean interface for the Domain Layer to retrieve and store domain entities without needing to understand the underlying SQL ⁴ dialects or table structures.




Figure 4.2. Internal Architecture

The internal architecture of the backend services, exemplified here in Figure 4.2 by the **Product service**, follows this strict layered pattern. The **product-api** component serves as the Presentation Layer, exposing REST endpoints to the client. This layer passes requests to the **product-domain**, which encapsulates the business logic and domain entities. The domain layer interacts with two downstream components: **product-data**, which handles persistence via the database, and **product-messaging**, which manages integrations with external banking systems. [28]

■ 4.1.4 Component Diagram

The component diagram illustrates the internal structure of the backend system and its interaction with the frontend, external bank APIs, and the database. It highlights how the conceptual layers described above are realized as actual software components.

³ JPA — Java Persistence API

⁴ SQL — Structured Query Language




Figure 4.3. Component Diagram

The **Frontend** communicates with the **Backend** through exposed REST APIs. Each request is routed to the appropriate service via dedicated controllers, which represent the entry point to the backend. The diagram shows the four main backend services: **Authorization Service**, **Bank Service**, **User Service**, and **Product Service**. Each service is isolated within its own boundary, emphasizing modularity and preventing business logic from bleeding across domains.

Within each service, the standard layered architecture is applied. The controller layer handles incoming HTTP requests and forwards them to the application API layer. The API layer defines service-specific interfaces and orchestrates calls to the domain layer.

Persistence responsibilities are delegated to repository and data components, which manage the physical database transactions. Communication with external banking systems is handled exclusively through adapter components. These APIs adapters provide a standardized interface for interacting with varying bank APIs, effectively shielding internal services from external changes, authentication complexities, and network failures.

Overall, the component diagram demonstrates a clean separation between the presentation, application, domain, and infrastructure layers. It highlights the loose coupling between services, clear responsibility boundaries, and the high extensibility of the system.

To provide further context on the system's modularity, the core business responsibilities delegated to each backend service are defined as follows:

- **Authorization Service** — acts as the dedicated security backbone and token vault for all external banking integrations. It centrally manages OAuth2 client registrations, secure token acquisition, and the entire cryptographic lifecycle. By handling the storage, retrieval, and automated refreshing of access and refresh tokens, it ensures that sensitive credentials are strictly isolated from the rest of the application's business logic.
- **User Service** — functions as the centralized user lifecycle manager. It completely isolates user identity and profile management from the financial domains.
- **Bank Connection Service** — manages the logical state of connected banking institutions. It abstracts the highly variable integration details of different Czech banks into a unified internal interface. Crucially, this service does not directly handle or store any security tokens. Instead, it delegates all token retrieval and refresh operations to the Authorization Service, allowing it to focus purely on orchestrating external bank communication and metadata.
- **Product Service** — serves as the core financial engine and data aggregation layer of the application. It is responsible for orchestrating financial data retrieval and normalizing data structures from various external banks into a single domain model.
- **Categorization Service** — is responsible for categorizing transactions and products based on the incoming information from the external bank API.

4.2 Data design

The data design defines how information is structured, stored, and processed within the application. It focuses on the consistency, integrity, and normalization of financial data retrieved from external APIs, as well as manually created fictive data. The design ensures that real and fictive accounts can coexist within a unified data model, providing a seamless experience for the user, regardless of the data source.

4.2.1 Database Model

The database model describes the relational structure used to persist application data, user profiles, and external integration metadata.

4.2.1.1 Authorized Client Table

OAuth2 authorized client	
PK	<u>client_registration_id</u> : varchar(100)
PK	<u>principal_name</u> : varchar(200)
	access_token_type: varchar(100)
	access_token_value: bytea
	access_token_issued_at: timestamp
	access_token_expires_at: timestamp
	access_token_scopes: varchar(1000)
	refresh_token_value: bytea
	refresh_token_issued_at: timestamp
	created_at: timestamp

Figure 4.4. Authorized Client Table

The OAuth2 authorized client table relies on a standardized schema utilized by Spring's OAuth2AuthorizedClientService. Its primary responsibility is to securely persist authorized client instances across application restarts. This includes storing the `access_token_value`, `refresh_token_value`, their respective expiration timestamps, and granted scopes. By delegating this to Spring Security, the application ensures that sensitive token lifecycle management is handled safely.

4.2.1.2 Bank Connection Table

Bank connection	
PK	<u>principal_name</u> : varchar(200)
PK	<u>client_registration_id</u> : varchar(100)
	bank_name: varchar(50)
	bank_connection_status: varchar(50)
	manually_created: boolean

Figure 4.5. Bank Connection Table

Bank connection table, detailed in Figure 4.5, stores application-specific data about a user's linked bank. It utilizes a composite primary key — `principal_name` (representing the user) and `client_registration_id` (representing the specific bank provider). It tracks the `bank_name`, the current `bank_connection_status`, and a `manually_created` flag.

4.2.1.3 Budget Table

Budget	
PK	<u>budget_id</u>: Long
	start_date: date
	amount: decimal(19, 4)
	currency: varchar(10)
	category: varchar(20)
	principal_name: String(100)

Figure 4.6. Budget Table

The **Budget** entity, presented in Figure 4.6, manages customer budgets. It is designed to store the financial limits users establish for their spending. The **start_date** field represents a recurring date when the budget is restarted. The defined financial limit is securely stored in the **amount** field. The currency code associated with this allocated amount, such as CZK or EUR, is specified within the **currency** field. The **category** attribute defines the specific spending group to which the limit applies. Finally, the **principal_name** attribute acts as a field that stores the unique identifier of the authenticated user who owns the budget.

4.2.1.4 Categorization Tables




Figure 4.7. Categorization Tables

The **Transaction category** table defines the standardized categories available within the system. It contains a primary key, **transaction_category_id**, and the **category_name** representing the classification (e.g., Groceries, Salary).

The **Transaction mapping** table is responsible for linking specific raw transaction information encountered in banking data (stored in the **transaction_name** field) to these standardized categories. It uses **transaction_mapping_id** as its primary key. The schema establishes a many-to-one relationship between **Transaction mapping** and **Transaction category**, allowing multiple distinct transaction names from various bank records to be cleanly grouped under a single unified category.

4.2.2.3 Product Service

The Product Service domain model represents the most data-intensive part of the application. It defines the relations between financial products, individual transactions, data aggregations, and user-defined budgets.




Figure 4.10. Domain Class Model — Product Service

The Product Service model forms the core financial logic of the system. It is centered around the **Product** entity, which serves as a unified representation for bank accounts.

Individual financial movements are recorded by the **Transaction** entity. This class captures core details such as booking dates, processing states (**TransactionStatus**), and cash flow direction (**TransactionDirection**). Complex counterpart details (debtor and creditor information) are extracted into a dedicated **RelatedParties** entity.

To directly support frontend analytics and dashboard visualizations, the model replaces specific time-based or category-based classes with versatile composition struc-

tures: **GroupedTransactions** and **GroupedProducts**. These entities logically organize raw financial records for efficient reporting. Additionally, the **AssetLiability** class connects to these grouped products to facilitate high-level financial summaries.

The **Budget** entity enables goal-oriented financial planning, allowing users to set monetary targets based on a specific category and track progress via the **BudgetStatus**. Finally, to ensure strict architectural consistency across the application, the classification enumerations utilized within this model, specifically **ProductCategory** and **TransactionCategory**, are located and maintained within the Shared Service module.

4.2.2.4 Shared Service

The Shared Service module functions as a foundational, cross-cutting component within the application architecture. Its primary objective is to centralize common programming logic, thereby eliminating code redundancy and ensuring strict systemic consistency across all independent functional domains.




Figure 4.11. Domain Class Model — Shared Service

The Shared Service model, illustrated in Figure 4.11, acts as a central model for the application's core enumerations. By isolating these classification structures into a dedicated module, the system ensures strict data consistency and eliminates code redundancy across all other functional domains.

The model defines three primary enumerations. The `TransactionCategory` enumeration provides a set of predefined labels used to automatically classify individual transactions. This structure includes everyday expense groups such as groceries, transportation, and entertainment, alongside broader classifications for income, fees, and uncategorized records.

Similarly, the `ProductCategory` enumeration standardizes the classification of all connected financial accounts. It encompasses standard banking products, including checking accounts and savings accounts, as well as complex financial instruments like investments, mortgages, and pensions.

Finally, the **ProductType** enumeration introduces a high-level financial grouping by categorizing products as an asset, a liability, or a neutral entity. This specific distinction

provides the fundamental logic required for the accurate calculation and visualization of the user's overall net worth within the consolidated dashboard.

Chapter 5

Implementation

This chapter details the practical realization of the personal finance system, translating the previously defined architectural concepts and data models into a fully functional application. The development process is structured into five logical phases, reflecting the requirements of building a secure and robust enterprise system. The chapter begins by detailing the realization of core application features; subsequent sections outline the security infrastructure, the application of caching and resilience patterns, and the containerized deployment strategy. Finally, the chapter describes the centralized error-handling logic.

5.1 Features

This section outlines the implementation of the primary user-facing functionalities. It details how the backend services process, aggregate, and deliver financial data to fulfill the core use cases of the application.

5.1.1 Dashboard

The Dashboard serves as the central entry point of the MojeFinance application, providing users with an immediate, high-level summary of their financial health. The primary objective of this interface is to aggregate complex, multi-source data from various connected banking institutions into a single, easily digestible view.




Figure 5.1. Feature — Dashboard

On this page, users can see financial visualizations that break down total assets and liabilities, allowing for a quick visual analysis of wealth distribution and debt across

different banks. The interface also includes a cash flow chart with all transactions from all connected accounts over the last six months, offering instant visibility into the latest cash flow activities without requiring navigation to the dedicated transactions module.

5.1.1.1 Cash Flow Summary

The cash flow summary visualizes financial movements from all linked accounts over the last six months. To support this feature efficiently, the backend **Product Service** utilizes the **TransactionsGroupingHelper** to systematically aggregate the data.

Transactions are grouped by their booking month and year, and then by category. Within each group, the system calculates the sum of all incoming funds and the sum of all outgoing expenses. Performing this aggregation on the server side significantly reduces the data payload transmitted over the network and completely offloads complex grouping calculations from the frontend application.

5.1.1.2 Assets and Liabilities

To provide an accurate calculation of the user's net worth, the system must logically distinguish between positive liquid balances and outstanding debts. The backend processes the user's entire financial portfolio by evaluating the **ProductCategory** enumeration assigned to each account.

Standard checking accounts, savings, and investments are automatically classified as assets, whereas credit cards, mortgages, and personal loans are categorized as liabilities. The service tier calculates the cumulative totals for both categories and groups products by their category. The resulting **AssetsAndLiabilitiesResponse** data structure allows the dashboard interface to clearly display the user's total available assets against their total liabilities. [29]

5.1.2 Banks Overview

The Banks Overview feature serves as the primary management interface for external financial integrations, exposing its lifecycle operations through the **BankController**.




Figure 5.2. Feature — Banks Overview

This interface allows users to view and manage their active links to supported financial institutions: Air Bank, Česká spořitelna, ČSOB, Komerční banka, and Raiffeisenbank. Users can initiate new connections from this view, which redirects them through

the secure consent flow of the respective bank to explicitly authorize data access. For existing connections, the interface displays current integration details and provides the capability to revoke access and disconnect an institution.

5.1.2.1 Data Retrieval

When the client sends GET request to the `/api/banks` endpoint, the `BankConnectionService` retrieves all registered integrations. It maps these internal entities to a standardized response, explicitly detailing their connection status (e.g., `CONNECTED` or `DISCONNECTED`) to indicate the validity of the underlying access tokens.

```
private void updateBankConnectionStatus(
    List<BankConnection> connectedBanks,
    String principalName
) {
    for (BankConnection connectedBank : connectedBanks) {
        String clientRegistrationId = connectedBank
            .getClientRegistrationId();

        boolean authorizedClientDoesNotExist =
            authorizedClientDoesNotExist(
                principalName,
                clientRegistrationId
            );
        if (authorizedClientDoesNotExist) {
            connectedBank.setBankConnectionStatus(
                BankConnectionStatus.DISCONNECTED
            );
            BankConnectionEntity bankEntityToUpdate =
                bankConnectionDomainMapper
                    .toBankConnectionEntity(connectedBank);
            bankConnectionRepository
                .updateConnectedBank(bankEntityToUpdate);
        }
    }
}
```

Listing 5.1. Update of Bank Connection Status

5.1.2.2 Connection Initialization

After the user authenticates on the target bank’s portal, the institution redirects them back to the application with a temporary authorization code.

The frontend transmits this code to the backend via a POST request to `/api/banks/connect`. The backend then delegates the token exchange to the Authorization Service and persists the active `BankConnection` entity in the database.

5.1.2.3 Disconnection

Users can terminate an existing bank connection at any time via a DELETE request to the `/api/banks/disconnect/bankId` endpoint.

The backend proactively interfaces with the token storage layer to permanently remove all associated OAuth2 materials. This deletion process ensures that the system retains no unauthorized access or refresh tokens linked to that specific institution.

5.1.3 Accounts Overview

The Accounts Overview feature provides a view of the user's financial products across all linked institutions. The client application retrieves this data by issuing a `GET` request to the `/api/products` endpoint. Then the `ProductService` retrieves all active bank connections associated with the user. For each valid connection, the service dynamically retrieves the required access tokens and executes real-time HTTP requests to the respective banking APIs.




Figure 5.3. Feature — Accounts Overview

This page displays individual account cards detailing the product name, associated banking institution, owner names, account identification details, and current available balance. Additionally, the interface includes intuitive client-side filtering tabs at the top, allowing users to isolate and view accounts belonging to a specific bank without triggering additional backend network requests.

To manage these external network communications, the backend uses Spring Cloud OpenFeign. By annotating methods with standard Spring web bindings, OpenFeign automatically generates the proxy implementations required to execute the requests.

Below is an example of the OpenFeign client used to interface with the Komerční banka API:

```

@FeignClient(
    name = "KBApiFeignClient",
    url = "${external.api.kb.base-url}",
    configuration = KBFeignConfig.class
)
public interface KBApiFeignClient {

    @GetMapping(value = "${external.api.kb.accounts-path}",
        consumes = MediaType.APPLICATION_JSON_VALUE)
    ResponseEntity<GetAccountListResponse> getAccounts(
        @RequestHeader("Authorization") String authorization
    );

    @GetMapping(value = "${external.api.kb.balances-path}",
        consumes = MediaType.APPLICATION_JSON_VALUE)
    ResponseEntity<GetAccountBalanceResponse> getAccountBalance(
        @RequestHeader("Authorization") String authorization,
        @PathVariable String id
    );

    @GetMapping(value = "${external.api.kb.transactions-path}",
        consumes = MediaType.APPLICATION_JSON_VALUE)
    ResponseEntity<GeAccountTransactionsResponse> getTransactions(
        @RequestHeader("Authorization") String authorization,
        @PathVariable String id,
        @RequestParam("fromDate") String fromDate,
        @RequestParam("toDate") String toDate
    );
}

```

Listing 5.2. Feign client for Komerční banka API integration

Once the raw data is retrieved, the backend must normalize the different JSON ¹ responses from the various institutions into a standardized **Product** data transfer object. This complex mapping process is implemented using the MapStruct library. MapStruct automatically generates high-performance, type-safe mapping code at compile time. [30]

¹ JSON — JavaScript Object Notation


```

@Mapper(componentModel = "spring")
public interface KBApiMapper extends TransactionsApiMapper {
    @Mapping(target = "products", source = "accounts")
    ProductsResponse toProductsResponse(GetAccountListResponse
    getAccountListResponse, @Context BankDetails bankDetails);

    @Mapping(target = "productId", source = "id")
    @Mapping(target = "productIdentification.iban", source =
    "identification.iban")
    @Mapping(target = "productIdentification.productNumber", source =
    "identification.other")
    @Mapping(target = "accountName", source = "nameI18N")
    @Mapping(target = "productName", source = "productI18N")
    @Mapping(target = "manuallyCreated", constant = "false")
    @Mapping(target = "bankCode", source = "servicer.bankCode")
    @Mapping(target = "bankDetails", expression = "java(bankDetails)")
    Product toDomainProduct(Account account, @Context BankDetails bankDe-
tails);

    //Other mapper methods
    ...
}

```

Listing 5.3. MapStruct Example

After normalization, the backend transmits the aggregated list to the client.

■ 5.1.4 Transactions Overview

The Transactions Overview feature provides a chronological view of the user's transaction history. The backend exposes a dedicated REST endpoint (`/api/products/clientRegistrationId/accountId/transactions`) that accepts temporal parameters (`fromDate` and `toDate`) to fetch specific transaction histories from the external banking APIs.




Figure 5.4. Feature — Transactions Overview

To optimize network performance and reduce client-side processing, the backend processes the data, filtering out pending transactions and grouping the completed ones hierarchically by their booking month and category. Raw data retrieval works the same way as described in the 5.1.3 section. The aggregated payload is then sent to the client.

The frontend application consumes this structured data to render the transactions view, enabling client-side filtering between income and expenses. The frontend uses this grouped payload to power a dynamic cash flow chart. As the user switches between different accounts or adjusts the date filters, this interactive chart updates in real-time, providing a visual summary of monthly spending habits alongside the detailed transaction list.

5.1.5 Budgeting

The Budgeting feature enables users to establish and monitor strict financial limits for specific transaction categories. The backend exposes a standard set of RESTful endpoints (`/api/budgets`) to handle the creation, modification, retrieval, and deletion of budgets.




Figure 5.5. Feature — Budgeting

The interface displays user budgets; each one contains a spending category with a user-defined financial limit and a budget restart date. Each budget card has a visual progress indicator, which dynamically updates to reflect the current amount spent versus the total allocated limit based on the system's categorized transaction data. Also, the interface has intuitive administrative controls that allow users to create new budgets, as well as edit or remove existing ones. To ensure user awareness, the interface emphasizes exceeded budgets by automatically sorting them to the top of the layout, applying visual warning tags, and triggering a global alert banner across the top of the view.

When creating a budget, the user can define a specific start date. This date acts as a recurring anchor; the backend automatically resets the budget cycle on this exact day every subsequent month. If the frontend payload omits this parameter, the system automatically injects the first day of the current month as the default fallback. The backend uses this timestamp to define the exact temporal window for evaluating the user's categorized expenses.

```

public void validateStartDate(Budget budget) {
    if (budget.getStartDate() == null) {
        LocalDate firstDayOfCurrentMonth = LocalDate.now()
            .withDayOfMonth(1);
        budget.setStartDate(firstDayOfCurrentMonth);
        log.info("Start date not provided. Setting start date to
            first day of current month: {}", firstDayOfCurrentMonth);
    }
}

public void updateBudgetStartDate(BudgetsResponse budgetsResponse) {
    for (Budget budget : budgetsResponse.getBudgets()) {
        YearMonth nextMonth = YearMonth
            .from(budget.getStartDate()).plusMonths(1);
        MonthDay monthDay = MonthDay.from(budget.getStartDate());
        LocalDate endDate = LocalDate.of(
            nextMonth.getYear(),
            nextMonth.getMonth(),
            monthDay.getDayOfMonth()
        );

        if (LocalDate.now().equals(endDate)) {
            log.info("Budget restarted for category: {}.
                Old start date: {}, New start date: {}",
                budget.getCategory(), budget.getStartDate(), endDate);
            budget.setStartDate(endDate);
            budget.setBudgetStatus(BudgetStatus.ACTIVE);
        }
    }
}

```

Listing 5.4. Start Date Processing

To provide accurate and real-time tracking, the backend service continuously evaluates the user's categorized expenses against these defined limits. If the calculated expenses for a specific category exceed the defined limit, the backend automatically updates the `budgetStatus` flag to `EXCEEDED`.

```

private void enrichBudgets(List<Budget> budgetsForDate, Map<String,
    BigDecimal> spentAmountPerCategory)
{
    budgetsForDate.forEach(budget -> {
        BigDecimal spentAmount = spentAmountPerCategory
            .getOrDefault(
                budget.getCategory().getDisplayName(),
                BigDecimal.ZERO
            );
        Amount spentAmountObj = Amount.builder()
            .value(spentAmount)
            .currency(budget.getAmount().getCurrency())
            .build();
        budget.setSpentAmount(spentAmountObj);

        if (spentAmount.compareTo(budget.getAmount().getValue()) >= 0) {
            budget.setBudgetStatus(BudgetStatus.EXCEEDED);
            log.info("Budget exceeded for category: {}.
                Spent amount: {}, Budget amount: {}",
                    budget.getCategory(), spentAmount,
                    budget.getAmount().getValue());
        } else {
            budget.setBudgetStatus(BudgetStatus.ACTIVE);
            log.info("Budget active for category: {}.
                Spent amount: {}, Budget amount: {}",
                    budget.getCategory(), spentAmount,
                    budget.getAmount().getValue());
        }

        log.info("Budget category: {}, Spent amount: {}",
            budget.getCategory(), spentAmount);
    });
}

```

Listing 5.5. Processing User's Categorized Expenses

■ 5.1.6 Categorization Service

The Categorization Service is a backend component that enriches raw financial data by assigning transactions and products to standardized categories.

Instead of relying on rigid, hard-coded rules or complex text matching, the implementation utilizes a highly efficient hybrid approach: a self-updating database dictionary paired with an integrated Google Gemini AI ² provider.

When the system receives data during synchronization, the service extracts the counterparty name and transaction direction from every transaction, and the product name from every product. The extracted names are queried against the `TransactionMappin-`

² AI — Artificial Intelligence

gEntity table in the database. If the system has seen this mapping before, it instantly retrieves the saved category.

Any names that are completely new and not found in the dictionary are sent to the Gemini AI API, which evaluates the names and returns the most probable categories from the system's predefined options. The newly acquired AI answers are saved back into the database mapping tables. This ensures the system continuously learns, which speeds up processing times and minimizes the need for external AI API calls. [31]

```

public Map<String, TransactionCategory> saveNewTransactionMappings
(Set<String> unmappedTransactions)
{
    Map<String, TransactionCategory> transactionMappings =
        new HashMap<>();

    for (String unmappedTransactionName : unmappedTransactions) {
        String prompt = buildPrompt(unmappedTransactionName);
        String geminiResponse = geminiProvider.askGemini(prompt);
        TransactionCategory mappedCategory = getTransactionCategory(
            unmappedTransactionName,
            geminiResponse
        );

        // Saving new mapping into the database
        ...

        transactionMappings
            .put(unmappedTransactionName, mappedCategory);
    }
    return transactionMappings;
}

@Override
public String askGemini(String prompt) {
    log.info("Asking Gemini with prompt: {}", prompt);

    GenerateContentResponse response;
    try {
        Client client = Client.builder()
            .apiKey(geminiApiKey)
            .build();

        response = getModels(client).generateContent(
            GEMINI_MODEL_NAME,
            prompt,
            null
        );
    } catch (Exception e) {
        log.error("Error while asking Gemini.", e);
        return EMPTY_STRING;
    }
    log.info("Received response from Gemini: {}", response.text());
    return Objects.requireNonNull(response.text()).trim();
}

```

Listing 5.6. AI-driven Categorization

5.1.7 Currency Exchange

When users connect bank accounts that hold foreign currencies (such as EUR or USD), directly summing these values with the local base currency (CZK) would result in mathematically incorrect totals for assets, liabilities, and cash flow charts.

To solve this, the backend implements a dedicated `CurrencyExchangeService`. This service intercepts foreign currency amounts and normalizes them into CZK before they are passed to the frontend or aggregated into dashboard totals.

```
public void exchangeProductBalances(List<Product> products) {
    for (Product product : products) {
        Amount originalAmount = product.getBalance();
        if (!CZK_CURRENCY_CODE.equals(originalAmount.getCurrency())) {
            Amount amountInCZK = currencyExchangeService
                .exchangeAmount(originalAmount);
            product.setBalance(amountInCZK);
        }
    }
}

@Override
public Amount exchangeAmount(Amount amount) {
    //Amount validation
    ...

    BigDecimal exchangeRate = externalApiProvider
        .getExchangeRates(amount.getCurrency());
    BigDecimal originalAmount = amount.getValue();
    BigDecimal result = originalAmount.multiply(exchangeRate);
    return Amount.builder()
        .value(result)
        .currency(CZK_CURRENCY_CODE)
        .build();
}
```

Listing 5.7. Currency Exchange

The system fetches daily exchange rate value using the Komerční banka Open API and caches it for 24 hours. Specifically, the service extracts the `cashBuy` rate for the requested foreign currency. This ensures the conversion reflects realistic banking market rates.

The conversion happens within the domain layer. The original foreign amount is multiplied by the fetched rate, generating a new `Amount` object with the CZK currency code.

5.2 Frontend Implementation

The client-side application is constructed using the React framework and bundled with Vite. It relies on standard ecosystem libraries, including React Router for client-side navigation and Axios for managing asynchronous HTTP communications.

To interact securely with the backend services, the frontend implements a centralized **AuthContext**. This context dictates the client-side authentication state and manages the Keycloak OAuth 2.0 lifecycle. It handles the initial login redirects, securely persists the access and refresh tokens within the browser's local storage, and orchestrates automatic, silent token rotation in the background before the short-lived access tokens expire.

Network communication is standardized through a custom Axios client configuration defined within **axiosClient.js**. Request interceptors automatically inject the current Bearer token into the authorization headers of all outgoing API calls. The application also defines dedicated routing components, such as the **BankCallback** page, to successfully intercept, parse, and process asynchronous redirection flows from external banking portals.

The user interface is divided into modular, domain-specific views, including the dashboard, transaction ledger, active bank connections, and budgeting interfaces. These components utilize custom React hooks (such as **useFetchConnectedBanks** and **useCashFlowAnalytics**) to consume the backend REST endpoints, aggregate the JSON payloads, and render real-time financial overviews without containing complex business logic themselves. To speed-up frontend development, the assistance of AI tools was employed, as the main goal of this project is the implementation of a robust backend layer.

5.3 Security

Securing user identities and protecting sensitive financial data are the foundational requirements of the implemented architecture. The system adopts a centralized, delegated authentication model to isolate security concerns from standard business logic.

5.3.1 Keycloak

Keycloak [32] is used as the centralized Identity and Access Management server. It handles the complete lifecycle of user authentication, independently managing user registrations, data, and securely storing credentials. By delegating these responsibilities entirely to Keycloak, the primary application backend is relieved of directly processing sensitive passwords and user data. The administrative configuration, including realm settings, token lifespans, and security policies, is provisioned through a predefined realm export file.

To facilitate the secure authorization process, the Keycloak realm defines two distinct clients corresponding to the application's architectural layers:

- **mojefinance-frontend** — This client is configured as a public client, designed for the user-facing web application. It is responsible for initiating the standard OAuth 2.0 authentication flow. When a user attempts to log in, the frontend client redirects them to Keycloak's hosted authentication pages. Upon successful credential verification, this client receives the resulting JSON Web Tokens, specifically the access token and refresh token, which are subsequently used to authorize requests to the server.
- **Backend client** — This client is configured as a bearer-only client, designed to protect the system's internal REST API endpoints. It does not actively initiate login flows. Instead, when the frontend transmits an HTTP request containing an access token in the authorization header, the backend client intercepts the transmission. It cryptographically verifies the token's signature using Keycloak's public keys, checks

the expiration status, and evaluates the embedded claims to enforce strict access control policies before executing any business logic.




Figure 5.6. Keycloak Administrator Dashboard

5.3.2 Authorization Service

The Authorization Service module encapsulates the complex OAuth 2.0 flow required for banking integration. It uses the Spring Security framework to manage mutual TLS (mTLS) communications, execute authorization code exchanges, automate token rotation, and securely persist client credentials. [33]

The core infrastructure is defined within the `OAuth2ClientManagerConfig` class. Institutions like Komerční banka and Česká spořitelna require specific mTLS certificates for secure communication. The configuration establishes specialized `RestTemplate` instances bound to explicit SSL ³ bundles for each institution.

```
public RestTemplate baseSslRestTemplate(String bundleName) {
    RestTemplate restTemplate = builder
        .setSslBundle(sslBundles.getBundle(bundleName))
        .build();

    restTemplate.setMessageConverters(getCustomMessageConverters());
    restTemplate
        .setErrorHandler(new OAuth2ErrorResponseErrorHandler());
    return restTemplate;
}
```

Listing 5.8. RestTemplate with SSL

These templates are then injected into customized `OAuth2AccessTokenResponseClient` implementations, ensuring that every outgoing token request utilizes the correct

³ SSL — Secure Sockets Layer

cryptographic identity and includes required proprietary headers, such as the specific API key headers.

```
@Bean
public OAuth2AccessTokenResponseClient
<OAuth2AuthorizationCodeGrantRequest>
authorizationCodeTokenResponseClient(
    @Qualifier("kbRestTemplate") RestTemplate kbTemplate,
    @Qualifier("csobRestTemplate") RestTemplate csobTemplate,
    @Qualifier("defaultRestTemplate") RestTemplate defaultTemplate)
{
    ...
}
```

Listing 5.9. Customized OAuth2AccessTokenResponseClient

When a user successfully completes the authentication on a bank’s external portal, the institution redirects the user back to the application along with a temporary authorization code. The `ExchangeTokenHelper` class intercepts this callback. It programmatically constructs an `OAuth2AuthorizationCodeGrantRequest`, executes the exchange to retrieve the initial access token and refresh token, and delegates the payload to the storage layer.

```
@Override
public void connectAuthorizedClient(ConnectAuthorizedClientRequest
request) throws ClientRegistrationNotFoundException {
    exchangeTokenHelper.exchangeToken(
        request.getClientRegistrationId(),
        request.getCode()
    );
}
```

Listing 5.10. Use of ExchangeTokenHelper

5.3.2.1 Token Storage and Encryption

Financial access tokens provide direct programmatic access to sensitive banking data and require strict protection at rest. Spring Security provides a standard `JdbcOAuth2AuthorizedClientService` to persist these OAuth2 entities within a relational database. To prevent credential exposure in the event of a database compromise, the application extends this capability by implementing a custom `AuthorizedClientService` decorator.

```

@Bean
public OAuth2AuthorizedClientService authorizedClientService(
    JdbcTemplate jdbcTemplate,
    ClientRegistrationRepository clientRegistrationRepository,
    @Value("${security.oauth2.encryption.password}") String password,
    @Value("${security.oauth2.encryption.salt}") String salt) {

    JdbcOAuth2AuthorizedClientService jdbcService =
        new JdbcOAuth2AuthorizedClientService(
            jdbcTemplate,
            clientRegistrationRepository
        );

    return new AuthorizedClientService(jdbcService, password, salt);
}

```

Listing 5.11. Token Storage

This custom service wraps the default JDBC ⁴ implementation. It intercepts the `OAuth2AuthorizedClient` object immediately before database persistence and applies symmetric encryption to the token strings. The encryption utilizes a configured secret password and salt via Spring Crypto's `TextEncryptor`. When the system later retrieves the client, the service automatically decrypts the values before returning them to the application logic.

```

@Override
public <T extends OAuth2AuthorizedClient> T loadAuthorizedClient(
    String clientRegistrationId, String principalName) {
    OAuth2AuthorizedClient authorizedClient = delegate.
        loadAuthorizedClient(clientRegistrationId, principalName);
    if (authorizedClient == null) {
        return null;
    }
    return (T) decrypt(authorizedClient);
}

@Override
public void saveAuthorizedClient(
    OAuth2AuthorizedClient authorizedClient, Authentication principal) {
    delegate.saveAuthorizedClient(encrypt(authorizedClient), principal);
}

```

Listing 5.12. Token Encryption and Decryption

⁴ JDBC — Java Database Connectivity

5.3.2.2 Token Rotation

Bank access tokens typically possess short lifespans, often expiring within short time frame. The system automates the token rotation process to maintain persistent bank connections without requiring repeated user intervention.

This automation is managed by the `DefaultOAuth2AuthorizedClientManager`, configured with an `OAuth2AuthorizedClientProvider`. When a business module requires an access token to fetch financial data, it invokes the `authorizeClient` method. The manager retrieves the encrypted token from the database, decrypts it, and evaluates its expiration timestamp.

```
@Bean
public OAuth2AuthorizedClientManager authorizedClientManager(
    ClientRegistrationRepository clientRegistrationRepository,
    OAuth2AuthorizedClientRepository authorizedClientRepository,
    OAuth2AccessTokenResponseClient<OAuth2RefreshTokenGrantRequest>
    refreshTokenTokenResponseClient) {

    OAuth2AuthorizedClientProvider authorizedClientProvider =
        OAuth2AuthorizedClientProviderBuilder.builder()
            .authorizationCode()
            .refreshToken(
                configurer -> configurer
                    .accessTokenResponseClient(
                        refreshTokenTokenResponseClient)
            )
            .build();

    DefaultOAuth2AuthorizedClientManager manager =
        new DefaultOAuth2AuthorizedClientManager(
            clientRegistrationRepository,
            authorizedClientRepository);
    manager.setAuthorizedClientProvider(authorizedClientProvider);
    return manager;
}
```

Listing 5.13. Token Rotation

If the access token is expired, the manager automatically halts the outgoing request. It constructs an `OAuth2RefreshTokenGrantRequest`, securely transmits the refresh token to the specific bank's authorization server using the pre-configured mTLS client, and acquires a new access token. The system then updates the encrypted database records with the new token values and returns the valid token to the calling service, allowing the data retrieval process to proceed without manual disruption.

5.4 Caching and Resilience

Integrating with multiple external banking APIs introduces inherent latency and reliability risks. Network timeouts, strict rate limits, or temporary outages on the provider's

side can severely degrade the user experience. To mitigate these risks and ensure high availability, the backend implements a dual strategy: data caching paired with fault-tolerance patterns using the Resilience4j library.

5.4.0.1 Caching Strategy

To minimize redundant network calls and accelerate data retrieval, the application utilizes Redis as a distributed cache.

The caching configuration is defined to allow different Time-To-Live policies depending on the volatility of the requested data. Standard banking data is cached for 30 minutes, whereas static daily metrics like exchange rates are cached for 24 hours.

```
@Configuration
@EnableCaching
public class RedisCacheConfig {

    @Bean
    public RedisCacheConfiguration cacheConfiguration() {
        return RedisCacheConfiguration.defaultCacheConfig()
            .entryTtl(Duration.ofMinutes(30))
            .disableCachingNullValues()
            // ... value serializer config
    }

    @Bean
    public RedisCacheManagerBuilderCustomizer
    redisCacheManagerBuilderCustomizer() {
        return (builder) -> builder
            .withCacheConfiguration("exchange-rates",
                RedisCacheConfiguration.defaultCacheConfig()
                    .entryTtl(Duration.ofHours(24))
                    .disableCachingNullValues()
                    // ... value serializer config
            )
    }
}
```

Listing 5.14. Redis Cache Configuration with variable TTLs

By annotating methods with `@Cacheable`, the system automatically intercepts requests and serves them directly from memory if the data is still valid.

```

@Override
//Resilience annotations
...
@Cacheable(value = "products", key = "#request.principalName + '-' +
#request.bankDetails.clientRegistrationId")
public ProductsResponse getProducts(ProductsMessagingRequest request) {
//Products retrieving logic
...
}

```

Listing 5.15. Cacheable Annotation Example

5.4.0.2 Retry and Circuit Breaker

When a cache miss occurs and an external API call must be made, the system protects itself by using Resilience4j's Retry and Circuit Breaker mechanisms.


```

@Override
@Retry(name = RESILIENCE_INSTANCE, fallbackMethod =
"fallbackGetProducts")
@CircuitBreaker(name = RESILIENCE_INSTANCE, fallbackMethod =
"fallbackGetProducts")
// Caching annotation
...
public ProductsResponse getProducts(ProductsMessagingRequest request) {
//Products retrieving logic
...
}

```

Listing 5.16. Retry and Circuit Breaker Annotations Example

- **Retry Mechanism:** Temporary network glitches are handled by automatically retrying failed requests. The system attempts the call up to 3 times with a 2-second delay between attempts before giving up.
- **Circuit Breaker:** If a bank API experiences a sustained outage, repeated failures will trip the Circuit Breaker. Once the failure rate reaches 50 percent (evaluated over a sliding window of 10 calls), the circuit opens. Further requests to that specific API are immediately blocked for 15 seconds to prevent overwhelming the failing service and to save local application threads. [34]

**Figure 5.7.** Circuit Breaker Pattern

This behavior is centrally defined in the application's configuration file:

```

resilience4j:
  circuitbreaker:
    instances:
      bankApi:
        slidingWindowSize: 10
        minimumNumberOfCalls: 5
        failureRateThreshold: 50
        waitDurationInOpenState: 15s
        permittedNumberOfCallsInHalfOpenState: 3
        automaticTransitionFromOpenToHalfOpenEnabled: true
  retry:
    instances:
      bankApi:
        maxAttempts: 3
        waitDuration: 2s
        retryExceptions:
          - cvut.fel.sit.shared.exception.ServiceException
          - feign.FeignException

```

Listing 5.17. Retry and Circuit Breaker configuration

When the circuit is open or retries are entirely exhausted, the system elegantly downgrades via explicitly defined fallback methods. Instead of throwing a critical HTTP 500 error to the frontend, these fallbacks intercept the failure and return safe, empty states.


```
private ProductsResponse fallbackGetProducts(
    ProductsMessagingRequest request,
    Throwable throwable
) {
    log.error("CircuitBreaker/Retry Fallback triggered for getProducts.
    Bank: {}. Reason: {}", request.getBankDetails().getBankName(),
    throwable.getMessage());
    return ProductsResponse.builder()
        .products(Collections.emptyList())
        .build();
}
```

Listing 5.18. Fallback Method Example

This ensures that a single failing bank API does not crash the entire synchronization process, guaranteeing that the dashboard remains functional and allows the user to view their other unaffected accounts seamlessly.


5.5 Deployment

This section explains how the application is built and prepared for release. First, it covers the Continuous Integration pipeline, which automatically compiles and tests the code whenever changes are made. Then, it describes how Docker is used to package the application into lightweight containers so that it can run on any server.

5.5.1 Continuous Integration Pipeline

To maintain code quality and ensure stability, the project integrates a Continuous Integration (CI) pipeline using GitLab CI. This automated workflow is triggered strictly on merge request events or direct commits to the default branch, preventing broken code from being integrated into the main repository. [35]

The pipeline runs within an isolated `maven:3.9-eclipse-temurin-21` environment. Upon successful compilation, the CI runner automatically extracts and stores the resulting `.jar` files as deployable artifacts, ready for the next stages of deployment.

**Figure 5.8.** Passed CI Pipeline

5.5.2 Docker

The backend application is containerized using a multi-stage Docker build to minimize the final image size. The initial build stage uses a heavier `maven:3.9.8-eclipse-temurin-21` base image to copy the project configuration files and source code, executing the Maven lifecycle to compile the application. The final runtime stage then starts from a lightweight `eclipse-temurin:21-jre-jammy` runtime, where it copies

the compiled `app.jar` from the builder stage and sets the entry point to execute the application.

To support local development, the system relies on a `docker-compose.yml` configuration that starts the entire architectural setup. The orchestrated services include a PostgreSQL database instance, a Redis container, a Keycloak server, and the MojeFinance backend application itself. The backend container employs the `depends_on` block to delay its startup sequence until both the PostgreSQL database and the Keycloak server successfully pass their internal health checks and evaluate to a `service_healthy` state. [36]

```
mojefinance-backend:
  build:
    context: .
    dockerfile: Dockerfile
  container_name: mojefinance-backend
  depends_on:
    postgres-db:
      condition: service_healthy
    keycloak:
      condition: service_healthy
  environment:
    SPRING_PROFILES_ACTIVE: docker
  ports:
    - "8081:8081"
  restart: unless-stopped
```

Listing 5.19. Docker Backend Setup

5.6 Error Handling

The application relies on a centralized approach to manage unexpected issues. At the core of this mechanism is a global exception handler, `GlobalExceptionHandler`, that monitors the entire application for errors. Using the Spring Boot `@ControllerAdvice` annotation, it intercepts exceptions thrown anywhere in the application before they reach the user. Once an error is caught, the handler processes it and returns a clean, standardized message to the frontend. This strategy ensures the application remains stable and the user receives clear feedback. [37]

To handle business logic errors, the application uses a custom `ServiceException`. It extends Java's `RuntimeException`, meaning that there is no need to write manual try-catch blocks for every potential error. This custom exception is thrown whenever an internal service encounters an unexpected state or fails to fetch data from an external banking API.

```
public class ServiceException extends RuntimeException {  
    public ServiceException(String message, Throwable cause) {  
        super(message, cause);  
    }  
  
    public ServiceException(String message) {  
        super(message);  
    }  
}
```

Listing 5.20. Service Exception

Chapter 6

Testing

This chapter outlines the testing strategy applied to the application. It details the methods and tools chosen to verify the correctness and reliability of the software. The testing process is divided into three main levels: evaluating isolated code components, checking the interactions between different system parts, and testing the entire integrated application from the perspective of a real user.

6.1 Unit Testing

The foundational layer of this strategy consists of unit tests, which evaluate individual application components in complete isolation. The backend uses JUnit 5 as the primary framework, paired with Mockito to generate mock objects for external dependencies. [38]

This approach ensures that core business logic, such as data normalization in mappers or financial calculations in domain helpers, functions correctly without requiring a live database or active network connections. By injecting mocked repositories and external API clients into the service layer, tests reliably trigger and verify specific edge cases. This includes checking custom error handling and fallback responses, ensuring the application behaves predictably under all conditions without interacting with real external banking systems.

```
class ProductHelperTest {
    @Mock
    private AuthorizationService authorizationService;

    @InjectMocks
    private ProductHelper productHelper;

    @BeforeEach
    void setUp() {
        MockitoAnnotations.openMocks(this);
    }

    //Individual tests
    ...
}
```

Listing 6.1. Mocking of ProductHelper Dependencies

The real `AuthorizationService` requires an active database and network connection. The test uses the `@Mock` annotation to create a fake, controlled version of the

service. The `@InjectMocks` annotation then automatically places this simulated dependency directly into the `ProductHelper` being tested. Finally, the `@BeforeEach` method ensures that a fresh set of mocks is initialized right before each individual test runs.

6.2 Integration Testing

After testing individual components, integration tests verify that different parts of the system work together correctly. They check if the application endpoints accept incoming requests, pass security checks, interact with the lower layers properly, and return the correct JSON responses. [39]

To demonstrate this, the following code snippet shows an integration test for the endpoint that fetches user products.

```

@Test
@WithMockUser(username = "testuser")
void getProducts_ShouldFetchFromExternalProviderAndMapToApiResponse()
throws Exception {
    when(externalApiProvider.getProducts(any()))
        .thenReturn(getProductsResponse());
    when(externalApiProvider.getAccountBalance(any()))
        .thenReturn(getProductBalance());
    when(categorizationService.categorizeProducts(any()))
        .thenReturn(getCategorizeProductsResponse());

    mockMvc.perform(get("/products")
        .header("Authorization", "Bearer test-token")
        .with(csrf())
        .contentType(MediaType.APPLICATION_JSON))
        .andExpect(status().isOk())
        .andExpect(jsonPath("$.products", hasSize(1)))
        .andExpect(jsonPath("$.products[0].productId",
            is("acc-9876")))
        .andExpect(jsonPath("$.products[0].productCategory",
            is("Loan")))
        .andExpect(jsonPath("$.products[0].accountName",
            is("accountName")))
        .andExpect(jsonPath("$.products[0].ownersNames",
            is(List.of("John Doe"))))
        .andExpect(jsonPath("$.products[0].productIdentification.iban",
            is("iban")))
        .andExpect(jsonPath("$.products[0].productIdentification
            .productNumber", is("productNumber")))
        .andExpect(jsonPath("$.products[0].balance.value",
            is(10.0)))
        .andExpect(jsonPath("$.products[0].balance.currency",
            is("CZK")))
        .andExpect(jsonPath("$.products[0].bankDetails
            .clientRegistrationId", is("kb")))
        .andExpect(jsonPath("$.products[0].bankDetails.bankName",
            is("KB")));
}

```

Listing 6.2. Get Products Integration Test

This test uses Spring's `MockMvc` tool to send a simulated HTTP GET request to the `/products` endpoint. First, it uses the `@WithMockUser` annotation to bypass the Keycloak security layer with a mocked authenticated user. Next, it provides controlled mock data for the external API connections and the internal categorization service. Finally, it executes the request and verifies that the application returns an HTTP 200 OK status, confirming that all JSON fields in the response exactly match the expected values and data types.

To ensure integration tests reflect a realistic environment, the application uses **Test-containers** for database interactions. This setup starts a real PostgreSQL instance inside a temporary Docker container during the test execution.

```
@Container
static PostgreSQLContainer<> postgres = new PostgreSQLContainer<>
("postgres:16-alpine");

@DynamicPropertySource
static void setProperties(DynamicPropertyRegistry registry) {
    registry.add("spring.datasource.url", postgres::getJdbcUrl);
    registry.add("spring.datasource.username", postgres::getUsername);
    registry.add("spring.datasource.password", postgres::getPassword);
}
```

Listing 6.3. Test Container Usage

The `@Container` annotation manages the lifecycle of the PostgreSQL container. Once the container is running, the `@DynamicPropertySource` method intercepts the generated connection details: database URL ¹, username, and password, and injects them directly into the Spring configuration. This approach tests the actual database queries and constraints, providing highly accurate validation of the data layer.

6.3 End-to-end Testing

The final layer of the testing strategy focuses on End-to-End testing. This approach evaluates the entire application flow, interacting with the frontend user interface as a real person would. Selenium WebDriver combined with JUnit 5 is used to automate browser actions.

Selenium is an open-source framework designed explicitly for web automation. It allows test scripts to directly control a real browser, simulating physical user actions like clicking buttons, typing in text fields, and navigating between pages. [40]

The system applies the Page Object Model design pattern. In this pattern, every screen of the application has a dedicated Java class that stores all the specific HTML ² locators and interactions for that page. This separates the test logic from the raw structural details of the web interface.

¹ URL — Uniform Resource Locator

² HTML — Hypertext Markup Language

```

public class DashboardPage {
    private final WebDriver driver;
    private final WebDriverWait wait;

    private final By loadingIndicator = By
        .cssSelector(".dashboard-grid.loading");
    private final By mainGrid = By
        .cssSelector(".dashboard-grid:not(.loading):not(.error)");
    private final By assetsTotalAmount = By
        .cssSelector(".assets-card .amount-large");

    // Constructor and other locators
    ...

    public boolean waitForDashboardToLoad() {
        try {
            wait.until(ExpectedConditions
                .invisibilityOfElementLocated(loadingIndicator));
            wait.until(ExpectedConditions
                .visibilityOfElementLocated(mainGrid));
            return true;
        } catch (TimeoutException e) {
            return false;
        }
    }

    public String getAssetsTotalAmount() {
        wait.until(ExpectedConditions
            .presenceOfElementLocated(assetsTotalAmount));
        return driver.findElement(assetsTotalAmount).getText().trim();
    }
}

```

Listing 6.4. Dashboard Page Object Model

The `DashboardPage` class begins by defining specific HTML locators using CSS ³ selectors. These `By` variables store the exact structural paths to visual elements on the web page.

The `waitForDashboardToLoad()` method pauses the test execution until the loading animation disappears and the main data grid becomes fully visible. This step prevents tests from failing due to slow network responses. Similarly, the `getAssetsTotalAmount()` method waits for the specific element to appear, extracts its text value, and returns it.

Building upon these page objects, the actual test classes define specific user scenarios. Because most features require an authenticated user, a common `TestBase` class handles the initial browser setup and the login sequence automatically before each test begins.

³ CSS — Cascading Style Sheets

Page Object	Test Scenario	Expected Outcome
Authentication	Successful User Login	The user submits valid credentials, and the application redirects to the Dashboard.
Dashboard	Core Component Visibility	The loading indicator hides, and the Cash Flow chart, Summary panel, Total Assets, and Total Liabilities cards appear.
Dashboard	Aggregated Balance Accuracy	Assets and Liabilities summaries show the expected total amounts.
Dashboard	Specific Account Balances	Specific target accounts appear and show their correct individual balances.
Accounts	Linked Products Render	Individual account cards load, displaying the product name, banking institution, account number, and balance.
Transactions	Transactions Displaying	Transactions are grouped by booking date and by category. Each transaction row displays data, including the counterparty, category, and amount.
Budgets	Budgets Displaying	All user-created budgets are present. Each budget card contains the spent amount, limit amount, category, and restarting date.
Budgets	New Budget Creation	It is possible to create a new valid budget.
Budgets	Budgets Managing	It is possible to edit or delete an existing budget.
Banks	Active Connections	The interface correctly lists all active financial institutions. Each connection card contains the name and connection status.
Banks	Creation of a New Connection	It is possible to connect a new bank from the list of supported financial institutions.
Banks	Bank Disconnection	It is possible to remove an existing connection to a bank.

Table 6.1. End-to-End Testing Scenarios



Chapter 7

Conclusion

The MojeFinance project was developed as a unified platform for personal financial management, demonstrating how Open Banking APIs can be employed to create a system for data aggregation. The core architectural objectives have been met, providing an environment for financial tracking. Beyond individual users, the platform could offer significant potential for financial advisors, who can use these aggregated insights to monitor client portfolios, track spending habits, and provide financial guidance.

However, this current iteration serves as a starting point rather than a final product. Future development should focus on maturing the system for real-world application, beginning with the transition from testing sandboxes to live production APIs. The platform's utility should be enhanced by introducing user management logic and the ability to add manual products and manual bank connections, ensuring users have a complete overview of their wealth. To further expand the platform's reach, upcoming versions can introduce support for business accounts, allowing business owners to track corporate cash flow alongside their personal finances within the same ecosystem.

The system's resilience and error handling must also be improved to provide a better user experience. This includes expanding the testing framework to better handle edge cases and refining the frontend to ensure seamless performance across all major browsers.



References

- [1] *You Need a Budget*.
<https://www.ynab.com/>.
- [2] *Monarch*.
<https://www.monarch.com/>.
- [3] *Pocket Guard*.
<https://pocketguard.com/>.
- [4] *Empower*.
<https://www.empower.com/>.
- [5] *Quicken Simplifi*.
<https://www.quicken.com/products/simplifi/>.
- [6] *Spendee*.
<https://www.spendee.com/>.
- [7] *BudgetBakers*.
<https://budgetbakers.com/>.
- [8] *PSD2*. 2026.
https://en.wikipedia.org/wiki/Payment_Services_Directive.
- [9] *Strong Customer Authentication*. 2024.
<https://auth0.com/blog/strong-customer-authentication-explained/>.
- [10] *KB API Developer Portal*.
<https://developers.kb.cz/>.
- [11] *Erste Developer Portal*.
<https://developers.erstegroup.com/>.
- [12] *Fio Internetbanking API*.
<https://www.fio.cz/bank-services/internetbanking-api>.
- [13] *Air Bank OpenAPI*.
<https://www.airbank.cz/open-api/open-api-account-info-v7-0.html##top>.
- [14] *Raiffeisenbank Developer Portal*.
<https://developers.rb.cz/premium/documentation>.
- [15] *ČSOB Developer Portal*.
<https://developers.rb.cz/premium/documentation>.
- [16] *Moneta Bank Developer Portal*.
<https://www.moneta.cz/zivnostnici-a-firmy/api-pro-vyvojare>.
- [17] *UniCredit Developer Portal*.
<https://developer.unicredit.eu/>.
- [18] *Spring Boot Framework*.
<https://www.ibm.com/think/topics/java-spring-boot>.

- [19] Joshua Bloch. *Effective Java. Second Edition*. 2008 .
- [20] *PostgreSQL*.
<https://www.postgresql.org/about/>.
- [21] *Redis Cache*. 2024 .
<https://medium.com/the-modern-scientist/caching-a-dive-into-in-memory-and-redis-caches-7b9491a1fa1b>.
- [22] *Resilience4j*. 2024 .
<https://www.j-labs.pl/en/tech-blog/fault-tolerance-with-resilience4j-and-spring-boot/>.
- [23] *Feign*.
https://cloud.spring.io/spring-cloud-netflix/multi/multi_spring-cloud-feign.html.
- [24] *Why use React*. 2025.
<https://adactio.medium.com/why-use-react-8ccd2ab4fea7>.
- [25] *Service Oriented Architecture*.
<https://aws.amazon.com/what-is/service-oriented-architecture/>.
- [26] *Microservices vs SOA*.
<https://aws.amazon.com/compare/the-difference-between-soa-microservices/>.
- [27] Neal Ford Mark Richards. *Fundamentals of Software Architecture*. O'Reilly, 2020 .
- [28] *Three Tier Architecture*.
<https://www.ibm.com/think/topics/three-tier-architecture>.
- [29] *Assets and Liabilities*. 2026.
<https://digit.business/insights/financial-literacy/what-is-an-asset-what-is-a-liability>.
- [30] *MapStruct*. 2025.
<https://www.baeldung.com/mapstruct>.
- [31] *Gemini API*.
<https://ai.google.dev/gemini-api/docs>.
- [32] *Keycloak*.
<https://www.keycloak.org/>.
- [33] *Spring Security Framework*. 2025.
<https://www.geeksforgeeks.org/advance-java/introduction-to-spring-security-and-its-features/>.
- [34] *Circuit Breaker*.
<https://resilience4j.readme.io/docs/circuitbreaker>.
- [35] *Continuous Integration*.
<https://aws.amazon.com/devops/continuous-integration/>.
- [36] *Docker Containerization*. 2026.
<https://www.geeksforgeeks.org/blogs/containerization-using-docker/>.
- [37] *Global Exception Handling*. 2025.
<https://medium.com/@AlexanderObregon/spring-boot-global-exception-handling-with-restcontrolleradvice-676c5b0b74ea>.
- [38] *Unit Testing*. 2026.
<https://www.geeksforgeeks.org/software-testing/unit-testing-software-testing/>.

- 
- [39] *Integration Testing*. 2025.
<https://www.ibm.com/think/topics/integration-testing>.
- [40] *Selenium*. 2023.
<https://ariangarshi.medium.com/getting-started-with-e2e-testing-using-selenium-webdriver-4b1074cae825>.



Appendix A

Zadání práce

Develop a web-based application to help private users manage their multiple personal banking accounts. The application should allow users to monitor accounts from various Czech financial institutions. It should aggregate multiple accounts into a unified dashboard, allow tracking of transactions from each account, and support budgeting. Research current solutions and design the basic functionalities of the proposed application. The developed application should include functionality for displaying bank accounts, transactions, and budgets. Identify a suitable framework for the frontend and backend of your application, and justify your selection of technologies. Create a set of scenarios to verify the functionality of both individual parts and the overall application. Based on the results, evaluate the application's current state and suggest additional functionality or other improvements.

Appendix B

Acronyms

PSD2	Payment Services Directive 2
API	Application Programming Interface
eIDAS	electronic IDentification, Authentication and trust Services
AIS	Account Information Services
PIS	Payment Initiation Services
SCA	Strong Customer Authentication
TLS	Transport Layer Security
OAuth	Open Authorization
QWAC	Qualified Website Authentication Certificate
AISP	Account Information Service Providers
FR	Functional Requirement
NFR	Non-functional Requirement
DOM	Document Object Model
UI	User Interface
SOA	Service-Oriented Architecture
REST	Representational State Transfer
HTTP	Java Persistence API
SQL	Structured Query Language
JSON	JavaScript Object Notation
AI	Artificial Intelligence
SSL	Secure Sockets Layer
JDBC	Java Database Connectivity
CI	Continuous Integration
URL	Uniform Resource Locator
HTML	Hypertext Markup Language
CSS	Cascading Style Sheets